

Recommender Systems

Given

- Users: U_1, \dots, U_n
- Movies: M_1, \dots, M_m
- Ratings: R_{ij}

Goal: Recommend movies to users

Challenges:

- Scale (millions of users, millions of movies)
- Cold Start (change in user base, change in content)
- Sparse Data (Not many users rank movies)

Neighborhood Methods

- (user, user) similarity measure
 - i.e. recommend same movies to similar users (requires info about users)
- (item, item) similarity measure
 - i.e. recommend movies that are similar (requires info about movies)

Pros:

- Intuitive / easy to explain
- No training
- Handles new users/items

Challenges:

- Users rate differently (bias)
- Ratings change over time (bias)

Feature Extraction - Content-Based:

Realistically:

- It's difficult to characterize movies and users with the right features
- Characterization of users and movies may not be accurate
 - If you are using genres for example, movies with varying degree of “comedy” will get the tag “comedy”.


Goal:

- Discover the best features in an automated way

Content-Based: assume you have features for movies - want to learn features for users

Collaborative filtering: want to learn features for both users and movies

$$P^{(j)} = \arg \min_P \frac{1}{\|M^{(j)}\|} \sum_{i \in M^{(j)}} (P^T Q^{(i)} - r_{ij})^2 + \lambda \|p\|^2$$



Regularization Term: a penalty on the size of the parameter p

Feature Extraction - Collaborative Filtering

Can't use SVD because R is sparse... BUT, we can formulate an optimization problem to solve:

$$\min_{p, q} \sum_{i, j \in R} (r_{ij} - p_i^T q_j)^2 + \lambda (\|p\|_F^2 + \|q\|_F^2)$$

$$R = PQ$$

To solve, take derivatives wrt P & Q . Then, just like Expectation-Maximization Algorithm from GMM:

1. Start with random Q
2. Get P
3. Improve Q
4. Repeat 2 & 3

Linear Regression

Motivation: Suppose we are given a curve $y = h(x)$, how can we evaluate whether it is a good fit to our data?

Compare $h(x_i)$ to y_i for all i .

Goal: For a given distance function d , find h where L is smallest.

$$L(h) = \sum_i d(h(x_i), y_i)$$

We want to minimize:

$$L(h) = \sum_i d(h(x_i), y_i)$$

We want to maximize:

$$\mathbf{L}(\mathbf{h}) = \mathbf{P}(\mathbf{Y} \mid \mathbf{h})$$

Assumptions Let's start by assuming our data was generated by a linear function plus some noise:

$$\vec{y} = h_{\beta}(X) + \vec{\epsilon}$$

$$h(x) = \beta_1 x \quad \checkmark$$

$$h(x) = \beta_0 + \beta_1 x \quad \checkmark$$

$$h(x) = \beta_0 + \beta_1 x + \beta_2 x^2 \quad \checkmark$$

$$h(x) = \beta_1 \log(x) + \beta_2 x^2 \quad \checkmark$$

$$h(x) = \beta_0 + \beta_1 x + \beta_1^2 x \quad \times$$

1. The relation between x (independent variable) and y (dependent variable) is linear in a parameter β .
2. ϵ_i are independent, identically distributed random variables following a $N(0, \sigma^2)$ distribution.
(Note: σ is constant)